

```

1  import numpy as np
2
3
4  class Softmax(object):
5
6      def __init__(self, dims=[10, 3073]):
7          self.init_weights(dims=dims)
8
9      def init_weights(self, dims):
10         """
11         Initializes the weight matrix of the Softmax classifier.
12         Note that it has shape (C, D) where C is the number of
13         classes and D is the feature size.
14         """
15         self.W = np.random.normal(size=dims) * 0.0001
16
17     def loss(self, X, y):
18         """
19         Calculates the softmax loss.
20
21         Inputs have dimension D, there are C classes, and we operate on minibatches
22         of N examples.
23
24         Inputs:
25         - X: A numpy array of shape (N, D) containing a minibatch of data.
26         - y: A numpy array of shape (N,) containing training labels; y[i] = c means
27           that X[i] has label c, where 0 <= c < C.
28
29         Returns a tuple of:
30         - loss as single float
31         """
32
33         # Initialize the loss to zero.
34         loss = 0.0
35
36         # ===== #
37         # YOUR CODE HERE:
38         #   Calculate the normalized softmax loss. Store it as the variable loss.
39         #   (That is, calculate the sum of the losses of all the training
40         #   set margins, and then normalize the loss by the number of
41         #   training examples.)
42         # ===== #
43         num_samples = X.shape[0]
44
45         all_scores = np.dot(X, self.W.T)
46         for i in range(num_samples):
47             sample_scores = all_scores[i] - np.max(all_scores[i])
48             sample_class_score = sample_scores[y[i]]
49             exp_sum = np.sum(np.exp(sample_scores))
50             loss = loss + np.log(exp_sum) - sample_class_score
51
52         loss = loss/num_samples
53
54         # ===== #
55         # END YOUR CODE HERE
56         # ===== #
57
58         return loss
59
60     def loss_and_grad(self, X, y):
61         """
62         Same as self.loss(X, y), except that it also returns the gradient.
63
64         Output: grad -- a matrix of the same dimensions as W containing
65           the gradient of the loss with respect to W.
66         """
67

```

```

68     # Initialize the loss and gradient to zero.
69     loss = 0.0
70     grad = np.zeros_like(self.W)
71
72     # ===== #
73     # YOUR CODE HERE:
74     # Calculate the softmax loss and the gradient. Store the gradient
75     # as the variable grad.
76     # ===== #
77     num_sample = X.shape[0]
78
79     all_scores = np.dot(X, self.W.T)
80     for i in range(num_sample):
81         sample_scores = all_scores[i] - np.max(all_scores[i])
82         sample_class_score = sample_scores[y[i]]
83
84         exp_sum = np.sum(np.exp(sample_scores))
85         loss = loss + np.log(exp_sum) - sample_class_score
86
87         sftmax = (np.exp(sample_scores) / exp_sum).reshape(-1, 1)
88         grad = grad + np.dot(sftmax, X[i].reshape(1, -1))
89         grad[y[i]] = grad[y[i]] - X[i]
90
91     loss = loss / num_sample
92     grad = grad / num_sample
93
94     # ===== #
95     # END YOUR CODE HERE
96     # ===== #
97
98     return loss, grad
99
100 def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
101     """
102     sample a few random elements and only return numerical
103     in these dimensions.
104     """
105
106     for i in np.arange(num_checks):
107         ix = tuple([np.random.randint(m) for m in self.W.shape])
108
109         oldval = self.W[ix]
110         self.W[ix] = oldval + h # increment by h
111         fxph = self.loss(X, y)
112         self.W[ix] = oldval - h # decrement by h
113         fxmh = self.loss(X, y) # evaluate f(x - h)
114         self.W[ix] = oldval # reset
115
116         grad_numerical = (fxph - fxmh) / (2 * h)
117         grad_analytic = your_grad[ix]
118         rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + abs(
119             grad_analytic))
120         print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical,
121             grad_analytic, rel_error))
122
123 def fast_loss_and_grad(self, X, y):
124     """
125     A vectorized implementation of loss_and_grad. It shares the same
126     inputs and outputs as loss_and_grad.
127     """
128
129     loss = 0.0
130     grad = np.zeros(self.W.shape) # initialize the gradient as zero
131
132     # ===== #
133     # YOUR CODE HERE:
134     # Calculate the softmax loss and gradient WITHOUT any for loops.
135     # ===== #

```

```

133     num_sample = X.shape[0]
134
135     all_scores = np.dot(X, self.W.T)
136     all_scores_stable = all_scores - np.max(all_scores, axis = 1, keepdims = True) #
sample x class
137
138     exp_sum = np.sum(np.exp(all_scores_stable), axis = 1, keepdims = True)
139     sftmax = np.exp(all_scores_stable) / exp_sum
140     sftmax = sftmax.clip(min = np.finfo(float).eps) # Added to avoid log0
141     loss = np.sum(-np.log(sftmax[np.arange(num_sample), y]))
142
143     sftmax[np.arange(num_sample), y] -= 1
144     grad = np.dot(sftmax.T, X)
145
146     loss = loss / num_sample
147     grad = grad / num_sample
148
149     # ===== #
150     # END YOUR CODE HERE
151     # ===== #
152
153     return loss, grad
154
155 def train(self, X, y, learning_rate=1e-3, num_iters=100,
156           batch_size=200, verbose=False):
157     """
158     Train this linear classifier using stochastic gradient descent.
159
160     Inputs:
161     - X: A numpy array of shape (N, D) containing training data; there are N
162         training samples each of dimension D.
163     - y: A numpy array of shape (N,) containing training labels; y[i] = c
164         means that X[i] has label 0 ≤ c < C for C classes.
165     - learning_rate: (float) learning rate for optimization.
166     - num_iters: (integer) number of steps to take when optimizing
167     - batch_size: (integer) number of training examples to use at each step.
168     - verbose: (boolean) If true, print progress during optimization.
169
170     Outputs:
171     A list containing the value of the loss function at each training iteration.
172     """
173     num_train, dim = X.shape
174     num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of
classes
175
176     self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of
self.W
177
178     # Run stochastic gradient descent to optimize W
179     loss_history = []
180
181     for it in np.arange(num_iters):
182         X_batch = None
183         y_batch = None
184
185         # ===== #
186         # YOUR CODE HERE:
187         # Sample batch_size elements from the training data for use in
188         # gradient descent. After sampling,
189         # - X_batch should have shape: (batch_size, dim)
190         # - y_batch should have shape: (batch_size,)
191         # The indices should be randomly generated to reduce correlations
192         # in the dataset. Use np.random.choice. It's okay to sample with
193         # replacement.
194         # ===== #
195         idx = np.random.choice(num_train, batch_size, replace = True)
196         X_batch = X[idx]

```

```

197     y_batch = y[idx]
198     # ===== #
199     # END YOUR CODE HERE
200     # ===== #
201
202     # evaluate loss and gradient
203     loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
204     loss_history.append(loss)
205
206     # ===== #
207     # YOUR CODE HERE:
208     #     Update the parameters, self.W, with a gradient step
209     # ===== #
210     self.W = self.W - (learning_rate * grad)
211
212     # ===== #
213     # END YOUR CODE HERE
214     # ===== #
215
216     if verbose and it % 100 == 0:
217         print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
218
219     return loss_history
220
221 def predict(self, X):
222     """
223     Inputs:
224     - X: N x D array of training data. Each row is a D-dimensional point.
225
226     Returns:
227     - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
228       array of length N, and each element is an integer giving the predicted
229       class.
230     """
231     y_pred = np.zeros(X.shape[1])
232     # ===== #
233     # YOUR CODE HERE:
234     #     Predict the labels given the training data.
235     # ===== #
236     all_scores = np.dot(X, self.W.T)
237     y_pred = np.argmax(all_scores, axis = 1)
238     # ===== #
239     # END YOUR CODE HERE
240     # ===== #
241
242     return y_pred
243
244

```