

```
import numpy as np
import matplotlib.pyplot as plt
```

```
class TwoLayerNet(object):
```

```
    """
```

A two-layer fully-connected neural network. The net has an input dimension of D , a hidden layer dimension of H , and performs classification over C classes. We train the network with a softmax loss function and L2 regularization on the weight matrices. The network uses a ReLU nonlinearity after the first fully connected layer.

In other words, the network has the following architecture:

input - fully connected layer - ReLU - fully connected layer - softmax

The outputs of the second fully-connected layer are the scores for each class.

```
    """
```

```
def __init__(self, input_size, hidden_size, output_size, std=1e-4):
```

```
    """
```

Initialize the model. Weights are initialized to small random values and biases are initialized to zero. Weights and biases are stored in the variable `self.params`, which is a dictionary with the following keys:

$W1$: First layer weights; has shape (H, D)

$b1$: First layer biases; has shape $(H,)$

$W2$: Second layer weights; has shape (C, H)

$b2$: Second layer biases; has shape $(C,)$

Inputs:

- `input_size`: The dimension D of the input data.

- `hidden_size`: The number of neurons H in the hidden layer.

- `output_size`: The number of classes C .

```
    """
```

```
self.params = {}
```

```
self.params['W1'] = std * np.random.randn(hidden_size, input_size)
```

```
self.params['b1'] = np.zeros(hidden_size)
```

```
self.params['W2'] = std * np.random.randn(output_size, hidden_size)
```

```
self.params['b2'] = np.zeros(output_size)
```

```
def loss(self, X, y=None, reg=0.0):
```

```
    """
```

Compute the loss and gradients for a two layer fully connected neural network.

Inputs:

- `X`: Input data of shape (N, D) . Each `X[i]` is a training sample.

- `y`: Vector of training labels. `y[i]` is the label for `X[i]`, and each `y[i]` is an integer in the range $0 \leq y[i] < C$. This parameter is optional; if it is not passed then we only return scores, and if it is passed then we instead return the loss and gradients.

- `reg`: Regularization strength.

Returns:

If `y` is `None`, return a matrix scores of shape (N, C) where `scores[i, c]` is the score for class `c` on input `X[i]`.

If `y` is not `None`, instead return a tuple of:

- `loss`: Loss (data loss and regularization loss) for this batch of training samples.

- `grads`: Dictionary mapping parameter names to gradients of those parameters with respect to the loss function; has the same keys as `self.params`.

```
    """
```

```
# Unpack variables from the params dictionary
```

```
W1, b1 = self.params['W1'], self.params['b1']
```

```

W2, b2 = self.params['W2'], self.params['b2']
N, D = X.shape

# Compute the forward pass
scores = None

# ===== #
# YOUR CODE HERE:
# Calculate the output scores of the neural network. The result
# should be (N, C). As stated in the description for this class,
# there should not be a ReLU layer after the second FC layer.
# The output of the second FC layer is the output scores. Do not
# use a for loop in your implementation.
# ===== #
softmax = lambda x: np.exp(x) / np.sum(np.exp(x), axis = 1, keepdims = True)
relu = lambda x: x * (x > 0)

h1 = relu(np.dot(X, W1.T) + b1)

scores = np.dot(h1, W2.T) + b2

# ===== #
# END YOUR CODE HERE
# ===== #

# If the targets are not given then jump out, we're done
if y is None:
    return scores

# Compute the loss
loss = None

# ===== #
# YOUR CODE HERE:
# Calculate the loss of the neural network. This includes the
# softmax loss and the L2 regularization for W1 and W2. Store the
# total loss in the variable loss. Multiply the regularization
# loss by 0.5 (in addition to the factor reg).
# ===== #

# scores is num_examples by num_classes
probabilities = softmax(scores)
y_hat = probabilities[np.arange(N), y]
L2_reg_term = 0.5 * reg * (np.sum(W1**2) + np.sum(W2**2))
loss = np.sum(-np.log(y_hat)) / N + L2_reg_term

# ===== #
# END YOUR CODE HERE
# ===== #

grads = {}

# ===== #
# YOUR CODE HERE:
# Implement the backward pass. Compute the derivatives of the
# weights and the biases. Store the results in the grads
# dictionary. e.g., grads['W1'] should store the gradient for
# W1, and be of the same size as W1.
# ===== #
softmax_grad = probabilities
softmax_grad[np.arange(N), y] -= 1
softmax_grad /= N

grads['W2'] = np.dot(softmax_grad.T, h1)
grads['b2'] = np.sum(softmax_grad, axis = 0)

```

```

grad_h2 = np.dot(softmax_grad,W2)

da = grad_h2
da[h1 <= 0] = 0 # relu
grads['W1'] = np.dot(da.T,X)
grads['b1'] = np.sum(da, axis = 0)

grads['W1'] += reg * W1
grads['W2'] += reg * W2

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

def train(self, X, y, X_val, y_val,
          learning_rate=1e-3, learning_rate_decay=0.95,
          reg=1e-5, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this neural network using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) giving training data.
    - y: A numpy array of shape (N,) giving training labels; y[i] = c means that
        X[i] has label c, where 0 <= c < C.
    - X_val: A numpy array of shape (N_val, D) giving validation data.
    - y_val: A numpy array of shape (N_val,) giving validation labels.
    - learning_rate: Scalar giving learning rate for optimization.
    - learning_rate_decay: Scalar giving factor used to decay the learning rate
        after each epoch.
    - reg: Scalar giving regularization strength.
    - num_iters: Number of steps to take when optimizing.
    - batch_size: Number of training examples to use per step.
    - verbose: boolean; if true print progress during optimization.
    """
    num_train = X.shape[0]
    iterations_per_epoch = max(num_train / batch_size, 1)

    # Use SGD to optimize the parameters in self.model
    loss_history = []
    train_acc_history = []
    val_acc_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        # ===== #
        # YOUR CODE HERE:
        # Create a minibatch by sampling batch_size samples randomly.
        # ===== #
        minibatch_idx = np.random.choice(np.arange(X.shape[0]), size=batch_size, replace=True)
        X_batch = X[minibatch_idx]
        y_batch = y[minibatch_idx]
        # ===== #
        # END YOUR CODE HERE
        # ===== #

        # Compute loss and gradients using the current minibatch
        loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
        loss_history.append(loss)

        # ===== #
        # YOUR CODE HERE:
        # Perform a gradient descent step using the minibatch to update

```

```

# all parameters (i.e., W1, W2, b1, and b2).
# ===== #

self.params["W1"] -= learning_rate*grads["W1"]
self.params["W2"] -= learning_rate*grads["W2"]
self.params["b1"] -= learning_rate*grads["b1"]
self.params["b2"] -= learning_rate*grads["b2"]

# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

# Every epoch, check train and val accuracy and decay learning rate.
if it % iterations_per_epoch == 0:
    # Check accuracy
    train_acc = (self.predict(X_batch) == y_batch).mean()
    val_acc = (self.predict(X_val) == y_val).mean()
    train_acc_history.append(train_acc)
    val_acc_history.append(val_acc)

    # Decay learning rate
    learning_rate *= learning_rate_decay

return {
    'loss_history': loss_history,
    'train_acc_history': train_acc_history,
    'val_acc_history': val_acc_history,
}

def predict(self, X):
    """
    Use the trained weights of this two-layer network to predict labels for
    data points. For each data point we predict scores for each of the C
    classes, and assign each data point to the class with the highest score.

    Inputs:
    - X: A numpy array of shape (N, D) giving N D-dimensional data points to
        classify.

    Returns:
    - y_pred: A numpy array of shape (N,) giving predicted labels for each of
        the elements of X. For all i, y_pred[i] = c means that X[i] is predicted
        to have class c, where 0 <= c < C.
    """
    y_pred = None

    # ===== #
    # YOUR CODE HERE:
    # Predict the class given the input data.
    # ===== #
    W1, b1 = self.params['W1'], self.params['b1']
    W2, b2 = self.params['W2'], self.params['b2']
    N, D = X.shape

    #softmax = lambda x: (np.exp(x)-np.max(x)) / np.sum((np.exp(x)- np.max(x)), axis = 1,
    keepdims = True)
    softmax = lambda x: np.exp(x) / np.sum(np.exp(x), axis = 1, keepdims = True)
    relu = lambda x: x * (x > 0)

    h1 = relu(np.dot(X,W1.T) + b1)
    scores = np.dot(h1,W2.T) + b2

    probabilities = softmax(scores)
    y_pred = np.argmax(probabilities,axis = 1)

```

```
# ===== #  
# END YOUR CODE HERE  
# ===== #
```

```
return y_pred
```