



Q1%

a) Receptive field of a neuron in  $CL_1$  is  $m_1 \times m_1$  or  $m_1$ .


Suppose  $m_1 = 2$ ,




Each neuron in  $CL_1$  is linked to  $m_1 \times m_1$  patch in input. Therefore, RF is  $m_1 \times m_1$ .



When  $m_1 = m_2 = 2$ , RF of a neuron in  $CL_2$  is  $3 \times 3$  or 3 which is  $m_1 + m_2 - 1$  since we deduct the overlap. Therefore, receptive field of each neuron in  $CL_2$  is  $m_1 + m_2 - 1 \times m_1 + m_2 - 1$  or  $m_1 + m_2 - 1$  when  $\text{stride}_1 = \text{stride}_2 = 1$ .



c)



It is not trivial to find receptive field, therefore one should consider a single layer.


Let  $r_e$  denote the number of features in feature map  $CL_e$  which contribute to generate one feature in  $CL_e$ .  $L$  is the last feature map.  $r_L = 1$

$k_e$ : kernel size ( $m_e$  in our case)

$s_e$ : stride

$r_{L-1} = k_L$ : we found that in a part.

Suppose we know  $r_e$  and we want to compute  $r_{e-1}$ .



$$k_e = 1$$


$$s_e = 3$$

$s_e r_e$  will cover all features that contribute,  
however, it will be covered  $s_e - 1$  more.

Therefore, formula will be

$$r_{e-1} = s_e r_e - (s_e - 1)$$

Suppose  $k_e > 1$ ,



$$k_e = 5$$


$$s_e = 3$$

Therefore, we will add  $k_e - 1$  features to cover all.

$$\begin{aligned} r_{e-1} &= s_e r_e - (s_e - 1) + k_e - 1 \\ &= s_e r_e + (k_e - s_e) \end{aligned}$$

or in our case


$$= s_e r_e + (m_e - s_e)$$



CL1

$$s_1 = 2$$

$$m_1 = 2$$



CL2

$$m_2 = 2$$

$$s_2 = ? = X$$

$$r_2 = 1$$

$$r_1 = s_2 + m_2 - s_e = \boxed{m_2} \Rightarrow r_{L-1} = m_L = 2 \checkmark$$

$$r_0 = s_1 r_1 + m_1 - s_1 = 2 * 2 + 2 - 2 = 4 \checkmark$$

Therefore, receptive field of  $CL_2$  is

$$r_2 = 1$$

$$r_1 = m_2$$

$$r_0 = s_1 m_2 + m_1 - s_1 \quad \text{or} \quad (s_1 m_2 + m_1 - s_1 \times s_1 m_2 + m_1 - s_1)$$

Same formula can be used for any layer, you just need to assign that layer to be  $L$ .

Therefore, receptive field of  $CL_1$  is

$$r_1 = r_L = 1$$

$$r_0 = r_{L-1} = m_0 = m_1 \quad \text{or} \quad (m_1 \times m_1)$$

Note that, receptive field of  $CL_1$  is independent of stride, therefore unchanged.

However, receptive field of  $CL_2$  can change.

$$\text{when } s_1 = 1 \Rightarrow RF_1 = m_2 + m_1 - 1 > 0$$

$$\text{when } s_1 = 2 \Rightarrow RF_2 = 2m_2 + m_1 - 2 > 0$$

$$\text{For minimum } m_2, m_1 \quad RF_1 = 1$$

$$RF_2 = 2m_2 + m_1 - 2 = 1$$

$$\text{For } m_2 = 2, m_1 = 1 \quad RF_1 = 2$$

$$RF_2 = 4 + 1 - 2 = 3 > 2$$

for  $m_2=1, m_1=2$   $RF_1 = 1$   
 $RF_2 = 1$

In result,  $LL_2$  receptive field gets bigger as  $s_1$  grows.

d) The recurrence relation is

$$r_{l-1} = s_l m_l + m_l - s_l \quad r_L = 1$$

which has the solution,

$$r_{L-1} = m_L$$

$$r_0 = \sum_{l=1}^L \left( (m_l - 1) \prod_{i=1}^{l-1} s_i \right) + 1$$

This was solved in practice problems, so I have only written the answer.

Note that we are asked for the  $k^{\text{th}}$  layer. therefore assign layer  $L$  as  $k^{\text{th}}$  layer. solution becomes

$$r_0 = \sum_{l=1}^K \left( (m_l - 1) \prod_{i=1}^{l-1} s_i \right) + 1$$

Solution makes sense when  $s_i = 1$ ,

$$r_0 = m_k + m_{k-1} + \dots + m_1 - k + 1$$

when  $k=2$

$$r_0 = m_2 - m_1 - 1 \text{ which I found in part b.}$$

Therefore, receptive layer of  $k^{\text{th}}$  layer is

$$\left( \sum_{l=1}^k \left( (m_{l-1})^{\prod_{i=1}^{l-1} s_i} \right) + 1 \times \sum_{l=1}^k \left( (m_{l-1})^{\prod_{i=1}^{l-1} s_i} \right) + 1 \right)$$

or

$$\sum_{l=1}^k \left( (m_{l-1})^{\prod_{i=1}^{l-1} s_i} \right) + 1$$

- e)
- a) Increasing filter size  $m_i$
  - b) Adding more layers
  - c) Increasing stride of filters except stride of last layer.

$$r_0 = \sum_{l=1}^k \left( (m_{l-1})^{\prod_{i=1}^{l-1} s_i} \right) + 1$$

$\nwarrow$  a)  $(m_{l-1}) \uparrow$   
 $\nwarrow$  b)  $k \uparrow$   
 $\nwarrow$  c)  $s_i \uparrow$

# CNN-Layers

February 26, 2023

## 0.1 Convolutional neural network layers

In this notebook, we will build the convolutional neural network layers. This will be followed by a spatial batchnorm, and then in the final notebook of this assignment, we will train a CNN to further improve the validation accuracy on CIFAR-10.

```
[1]: ## Import and setups

import time

import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from utils.data_utils import get_CIFAR10_data
from utils.gradient_check import eval_numerical_gradient,   
    ↪eval_numerical_gradient_array
from utils.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
    ↪autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

## 0.2 Implementing CNN layers

Just as we implemented modular layers for fully connected networks, batch normalization, and dropout, we'll want to implement modular layers for convolutional neural networks. These layers are in `nndl/conv_layers.py`.

### 0.2.1 Convolutional forward pass

Begin by implementing a naive version of the forward pass of the CNN that uses `for` loops. This function is `conv_forward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a triple `for` loop.

After you implement `conv_forward_naive`, test your implementation by running the cell below.

```
[2]: x_shape = (2, 3, 4, 4)
w_shape = (3, 3, 4, 4)
x = np.linspace(-0.1, 0.5, num=np.prod(x_shape)).reshape(x_shape)
w = np.linspace(-0.2, 0.3, num=np.prod(w_shape)).reshape(w_shape)
b = np.linspace(-0.1, 0.2, num=3)

conv_param = {'stride': 2, 'pad': 1}
out, _ = conv_forward_naive(x, w, b, conv_param)
correct_out = np.array([[[[-0.08759809, -0.10987781],
                           [-0.18387192, -0.2109216 ]],
                          [[ 0.21027089,  0.21661097],
                           [ 0.22847626,  0.23004637]],
                          [[ 0.50813986,  0.54309974],
                           [ 0.64082444,  0.67101435]]],
                         [[[-0.98053589, -1.03143541],
                           [-1.19128892, -1.24695841]],
                          [[ 0.69108355,  0.66880383],
                           [ 0.59480972,  0.56776003]],
                          [[ 2.36270298,  2.36904306],
                           [ 2.38090835,  2.38247847]]]])

# Compare your output to ours; difference should be around 1e-8
print('Testing conv_forward_naive')
print('difference: ', rel_error(out, correct_out))
```

```
Testing conv_forward_naive
difference:  2.2121476417505994e-08
```

### 0.2.2 Convolutional backward pass

Now, implement a naive version of the backward pass of the CNN. The function is `conv_backward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a quadruple `for` loop.

After you implement `conv_backward_naive`, test your implementation by running the cell below.

```
[3]: x = np.random.randn(4, 3, 5, 5)
w = np.random.randn(2, 3, 3, 3)
b = np.random.randn(2,)
dout = np.random.randn(4, 2, 5, 5)
```



```

conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_forward_naive(x,w,b,conv_param)

dx_num = eval_numerical_gradient_array(lambda x: conv_forward_naive(x, w, b,
    ↪conv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_forward_naive(x, w, b,
    ↪conv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_forward_naive(x, w, b,
    ↪conv_param)[0], b, dout)

out, cache = conv_forward_naive(x, w, b, conv_param)
dx, dw, db = conv_backward_naive(dout, cache)

# Your errors should be around 1e-9'
print('Testing conv_backward_naive function')
print('dx error: ', rel_error(dx, dx_num))
print('dw error: ', rel_error(dw, dw_num))
print('db error: ', rel_error(db, db_num))

```

```

Testing conv_backward_naive function
dx error:  1.916590455692754e-08
dw error:  7.4069185122532045e-09
db error:  1.8782924169840352e-11

```

### 0.2.3 Max pool forward pass

In this section, we will implement the forward pass of the max pool. The function is `max_pool_forward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_forward_naive`, test your implementation by running the cell below.

```

[4]: x_shape = (2, 3, 4, 4)
x = np.linspace(-0.3, 0.4, num=np.prod(x_shape)).reshape(x_shape)
pool_param = {'pool_width': 2, 'pool_height': 2, 'stride': 2}

out, _ = max_pool_forward_naive(x, pool_param)

correct_out = np.array([[[[-0.26315789, -0.24842105],
                           [-0.20421053, -0.18947368]],
                          [[-0.14526316, -0.13052632],
                           [-0.08631579, -0.07157895]],
                          [[-0.02736842, -0.01263158],
                           [ 0.03157895,  0.04631579]]],
                        [[[ 0.09052632,  0.10526316],
                           [ 0.14947368,  0.16421053]]],

```

```

[[ 0.20842105,  0.22315789],
 [ 0.26736842,  0.28210526]],
 [[ 0.32631579,  0.34105263],
 [ 0.38526316,  0.4         ]]]])

# Compare your output with ours. Difference should be around 1e-8.
print('Testing max_pool_forward_naive function:')
print('difference: ', rel_error(out, correct_out))

```

Testing max\_pool\_forward\_naive function:  
 difference: 4.1666665157267834e-08

### 0.2.4 Max pool backward pass

In this section, you will implement the backward pass of the max pool. The function is `max_pool_backward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_backward_naive`, test your implementation by running the cell below.

```

[5]: x = np.random.randn(3, 2, 8, 8)
dout = np.random.randn(3, 2, 4, 4)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

dx_num = eval_numerical_gradient_array(lambda x: max_pool_forward_naive(x,
    ↪pool_param)[0], x, dout)

out, cache = max_pool_forward_naive(x, pool_param)
dx = max_pool_backward_naive(dout, cache)

# Your error should be around 1e-12
print('Testing max_pool_backward_naive function:')
print('dx error: ', rel_error(dx, dx_num))

```

Testing max\_pool\_backward\_naive function:  
 dx error: 3.275620431353226e-12

## 0.3 Fast implementation of the CNN layers

Implementing fast versions of the CNN layers can be difficult. We will provide you with the fast layers implemented by `utils`. They are provided in `utils/fast_layers.py`.

The fast convolution implementation depends on a Cython extension ('pip install Cython' to your virtual environment); to compile it you need to run the following from the `utils` directory:

```
python setup.py build_ext --inplace
```

**NOTE:** The fast implementation for pooling will only perform optimally if the pooling regions are non-overlapping and tile the input. If these conditions are not met then the fast pooling implementation will not be much faster than the naive implementation.

You can compare the performance of the naive and fast versions of these layers by running the cell below.

You should see pretty drastic speedups in the implementation of these layers. On our machine, the forward pass speeds up by 17x and the backward pass speeds up by 840x. Of course, these numbers will vary from machine to machine, as well as on your precise implementation of the naive layers.

```
[6]: from utils.fast_layers import conv_forward_fast, conv_backward_fast
    from time import time

    x = np.random.randn(100, 3, 31, 31)
    w = np.random.randn(25, 3, 3, 3)
    b = np.random.randn(25,)
    dout = np.random.randn(100, 25, 16, 16)
    conv_param = {'stride': 2, 'pad': 1}

    t0 = time()
    out_naive, cache_naive = conv_forward_naive(x, w, b, conv_param)
    t1 = time()
    out_fast, cache_fast = conv_forward_fast(x, w, b, conv_param)
    t2 = time()

    print('Testing conv_forward_fast:')
    print('Naive: %fs' % (t1 - t0))
    print('Fast: %fs' % (t2 - t1))
    print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
    print('Difference: ', rel_error(out_naive, out_fast))

    t0 = time()
    dx_naive, dw_naive, db_naive = conv_backward_naive(dout, cache_naive)
    t1 = time()
    dx_fast, dw_fast, db_fast = conv_backward_fast(dout, cache_fast)
    t2 = time()

    print('\nTesting conv_backward_fast:')
    print('Naive: %fs' % (t1 - t0))
    print('Fast: %fs' % (t2 - t1))
    print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
    print('dx difference: ', rel_error(dx_naive, dx_fast))
    print('dw difference: ', rel_error(dw_naive, dw_fast))
    print('db difference: ', rel_error(db_naive, db_fast))
```

Testing conv\_forward\_fast:

Naive: 2.644079s

Fast: 0.007051s

Speedup: 375.006695x

Difference: 1.0760953764912739e-10

Testing conv\_backward\_fast:

Naive: 4.269495s  
Fast: 0.007568s  
Speedup: 564.177625x  
dx difference: 1.0012070746726426e-10  
dw difference: 1.8655646296229697e-12  
db difference: 3.9094552906934445e-15

```
[7]: from utils.fast_layers import max_pool_forward_fast, max_pool_backward_fast

x = np.random.randn(100, 3, 32, 32)
dout = np.random.randn(100, 3, 16, 16)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

t0 = time()
out_naive, cache_naive = max_pool_forward_naive(x, pool_param)
t1 = time()
out_fast, cache_fast = max_pool_forward_fast(x, pool_param)
t2 = time()

print('Testing pool_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('fast: %fs' % (t2 - t1))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive = max_pool_backward_naive(dout, cache_naive)
t1 = time()
dx_fast = max_pool_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting pool_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('dx difference: ', rel_error(dx_naive, dx_fast))
```

Testing pool\_forward\_fast:  
Naive: 0.213794s  
fast: 0.004436s  
speedup: 48.192562x  
difference: 0.0

Testing pool\_backward\_fast:  
Naive: 0.581855s  
speedup: 72.587877x  
dx difference: 0.0

## 0.4 Implementation of cascaded layers

We've provided the following functions in `nndl/conv_layer_utils.py`: - `conv_relu_forward` - `conv_relu_backward` - `conv_relu_pool_forward` - `conv_relu_pool_backward`

These use the fast implementations of the conv net layers. You can test them below:

```
[8]: from nndl.conv_layer_utils import conv_relu_pool_forward, \
      ↪conv_relu_pool_backward

x = np.random.randn(2, 3, 16, 16)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

out, cache = conv_relu_pool_forward(x, w, b, conv_param, pool_param)
dx, dw, db = conv_relu_pool_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_pool_forward(x, w, \
      ↪b, conv_param, pool_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_pool_forward(x, w, \
      ↪b, conv_param, pool_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_pool_forward(x, w, \
      ↪b, conv_param, pool_param)[0], b, dout)

print('Testing conv_relu_pool')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

Testing conv\_relu\_pool

dx error: 7.939299810561729e-09

dw error: 6.562297642758617e-09

db error: 3.100564107304263e-11

```
[9]: from nndl.conv_layer_utils import conv_relu_forward, conv_relu_backward

x = np.random.randn(2, 3, 8, 8)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_relu_forward(x, w, b, conv_param)
dx, dw, db = conv_relu_backward(dout, cache)
```

```

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_forward(x, w, b,
    ↪conv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_forward(x, w, b,
    ↪conv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_forward(x, w, b,
    ↪conv_param)[0], b, dout)

print('Testing conv_relu:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))

```

Testing conv\_relu:

```

dx error:  4.004510301971303e-09
dw error:  1.5116437241970764e-09
db error:  2.7145990866722925e-11

```

## 0.5 What next?

We saw how helpful batch normalization was for training FC nets. In the next notebook, we'll implement a batch normalization for convolutional neural networks, and then finish off by implementing a CNN to improve our validation accuracy on CIFAR-10.

# CNN-BatchNorm

February 26, 2023

## 0.1 Spatial batch normalization

In fully connected networks, we performed batch normalization on the activations. To do something equivalent on CNNs, we modify batch normalization slightly.

Normally batch-normalization accepts inputs of shape  $(N, D)$  and produces outputs of shape  $(N, D)$ , where we normalize across the minibatch dimension  $N$ . For data coming from convolutional layers, batch normalization accepts inputs of shape  $(N, C, H, W)$  and produces outputs of shape  $(N, C, H, W)$  where the  $N$  dimension gives the minibatch size and the  $(H, W)$  dimensions give the spatial size of the feature map.

How do we calculate the spatial averages? First, notice that for the  $C$  feature maps we have (i.e., the layer has  $C$  filters) that each of these ought to have its own batch norm statistics, since each feature map may be picking out very different features in the images. However, within a feature map, we may assume that across all inputs and across all locations in the feature map, there ought to be relatively similar first and second order statistics. Hence, one way to think of spatial batch-normalization is to reshape the  $(N, C, H, W)$  array as an  $(N*H*W, C)$  array and perform batch normalization on this array.

Since spatial batch norm and batch normalization are similar, it'd be good to at this point also copy and paste our prior implemented layers from HW #4. Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4: - layers.py for your FC network layers, as well as batchnorm and dropout. - layer\_utils.py for your combined FC network layers. - optim.py for your optimizers.

Be sure to place these in the `nnd1/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

If you use your prior implementations of the batchnorm, then your spatial batchnorm implementation may be very short. Our implementations of the forward and backward pass are each 6 lines of code.

```
[1]: ## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
```

```

from nndl.conv_layers import *
from utils.data_utils import get_CIFAR10_data
from utils.gradient_check import eval_numerical_gradient, \
    ↪eval_numerical_gradient_array
from utils.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
    ↪autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

```

## 0.2 Spatial batch normalization forward pass

Implement the forward pass, `spatial_batchnorm_forward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

```

[2]: # Check the training-time forward pass by checking means and variances
      # of features both before and after spatial batch normalization

N, C, H, W = 2, 3, 4, 5
x = 4 * np.random.randn(N, C, H, W) + 10

print('Before spatial batch normalization:')
print('  Shape: ', x.shape)
print('  Means: ', x.mean(axis=(0, 2, 3)))
print('  Stds: ', x.std(axis=(0, 2, 3)))

# Means should be close to zero and stds close to one
gamma, beta = np.ones(C), np.zeros(C)
bn_param = {'mode': 'train'}
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization:')
print('  Shape: ', out.shape)
print('  Means: ', out.mean(axis=(0, 2, 3)))
print('  Stds: ', out.std(axis=(0, 2, 3)))

# Means should be close to beta and stds close to gamma

```



```

gamma, beta = np.asarray([3, 4, 5]), np.asarray([6, 7, 8])
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization (nontrivial gamma, beta):')
print('  Shape: ', out.shape)
print('  Means: ', out.mean(axis=(0, 2, 3)))
print('  Stds: ', out.std(axis=(0, 2, 3)))

```

Before spatial batch normalization:

```

Shape: (2, 3, 4, 5)
Means: [ 9.77859078  9.04656259 10.2300046 ]
Stds:  [4.68708021 4.17886643 4.67689384]

```

After spatial batch normalization:

```

Shape: (2, 3, 4, 5)
Means: [ 5.55111512e-18  6.78623824e-16 -4.44089210e-17]
Stds:  [0.99999964 0.99999977 0.9999998 ]

```

After spatial batch normalization (nontrivial gamma, beta):

```

Shape: (2, 3, 4, 5)
Means: [6. 7. 8.]
Stds:  [2.99999891 3.99999909 4.99999902]

```

### 0.3 Spatial batch normalization backward pass

Implement the backward pass, `spatial_batchnorm_backward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

```

[3]: N, C, H, W = 2, 3, 4, 5
x = 5 * np.random.randn(N, C, H, W) + 12
gamma = np.random.randn(C)
beta = np.random.randn(C)
dout = np.random.randn(N, C, H, W)

bn_param = {'mode': 'train'}
fx = lambda x: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
fg = lambda a: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
fb = lambda b: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]

dx_num = eval_numerical_gradient_array(fx, x, dout)
da_num = eval_numerical_gradient_array(fg, gamma, dout)
db_num = eval_numerical_gradient_array(fb, beta, dout)

_, cache = spatial_batchnorm_forward(x, gamma, beta, bn_param)
dx, dgamma, dbeta = spatial_batchnorm_backward(dout, cache)
print('dx error: ', rel_error(dx_num, dx))
print('dgamma error: ', rel_error(da_num, dgamma))
print('dbeta error: ', rel_error(db_num, dbeta))

```

```

dx error:  5.586411465744636e-08
dgamma error:  1.7706480782599564e-12
dbeta error:  3.2917466860558243e-12

```

[ ]:

# CNN

February 26, 2023

## 1 Convolutional neural networks

In this notebook, we'll put together our convolutional layers to implement a 3-layer CNN. Then, we'll ask you to implement a CNN that can achieve  $> 65\%$  validation error on CIFAR-10.

If you have not completed the Spatial BatchNorm Notebook, please see the following description from that notebook:

Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their layer implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4: - layers.py for your FC network layers, as well as batchnorm and dropout. - layer\_utils.py for your combined FC network layers. - optim.py for your optimizers.

Be sure to place these in the `nndl/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

```
[1]: # As usual, a bit of setup

import numpy as np
import matplotlib.pyplot as plt
from nndl.cnn import *
from utils.data_utils import get_CIFAR10_data
from utils.gradient_check import eval_numerical_gradient_array, \
    eval_numerical_gradient
from nndl.layers import *
from nndl.conv_layers import *
from utils.fast_layers import *
from utils.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
```

```
# see http://stackoverflow.com/questions/1907993/
↳ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

```
[2]: # Load the (preprocessed) CIFAR10 data.

data = get_CIFAR10_data()
for k in data.keys():
    print('{}: {}'.format(k, data[k].shape))
```

```
X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

## 1.1 Three layer CNN

In this notebook, you will implement a three layer CNN. The `ThreeLayerConvNet` class is in `mndl/cnn.py`. You'll need to modify that code for this section, including the initialization, as well as the calculation of the loss and gradients. You should be able to use the building blocks you have either earlier coded or that we have provided. Be sure to use the fast layers.

The architecture of this CNN will be:

conv - relu - 2x2 max pool - affine - relu - affine - softmax

We won't use batchnorm yet. You've also done enough of these to know how to debug; use the cells below.

Note: As we are implementing several layers CNN networks. The gradient error can be expected for the `eval_numerical_gradient()` function. If your `W1` max relative error and `W2` max relative error are around or below 0.01, they should be acceptable. Other errors should be less than  $1e-5$ .

```
[3]: num_inputs = 2
input_dim = (3, 16, 16)
reg = 0.0
num_classes = 10
X = np.random.randn(num_inputs, *input_dim)
y = np.random.randint(num_classes, size=num_inputs)

model = ThreeLayerConvNet(num_filters=3, filter_size=3,
                           input_dim=input_dim, hidden_dim=7,
```

```

dtype=np.float64)
loss, grads = model.loss(X, y)
for param_name in sorted(grads):
    f = lambda _: model.loss(X, y)[0]
    param_grad_num = eval_numerical_gradient(f, model.params[param_name],
↪verbose=False, h=1e-6)
    e = rel_error(param_grad_num, grads[param_name])
    print('{} max relative error: {}'.format(param_name,
↪rel_error(param_grad_num, grads[param_name])))

```

```

W1 max relative error: 0.00014266947341464817
W2 max relative error: 0.010526612646796649
W3 max relative error: 0.0002806485855424487
b1 max relative error: 7.116381199585126e-05
b2 max relative error: 6.593377672518949e-07
b3 max relative error: 1.4712289610079897e-09

```

### 1.1.1 Overfit small dataset

To check your CNN implementation, let's overfit a small dataset.

```

[4]: num_train = 100
small_data = {
    'X_train': data['X_train'][:num_train],
    'y_train': data['y_train'][:num_train],
    'X_val': data['X_val'],
    'y_val': data['y_val'],
}

model = ThreeLayerConvNet(weight_scale=1e-2)

solver = Solver(model, small_data,
                num_epochs=10, batch_size=50,
                update_rule='adam',
                optim_config={
                    'learning_rate': 1e-3,
                },
                verbose=True, print_every=1)
solver.train()

```

```


(Iteration 1 / 20) loss: 2.282205
(Epoch 0 / 10) train acc: 0.300000; val_acc: 0.124000
(Iteration 2 / 20) loss: 4.254238
(Epoch 1 / 10) train acc: 0.240000; val_acc: 0.122000
(Iteration 3 / 20) loss: 2.792451
(Iteration 4 / 20) loss: 2.095923
(Epoch 2 / 10) train acc: 0.370000; val_acc: 0.166000
(Iteration 5 / 20) loss: 2.132012
(Iteration 6 / 20) loss: 1.821241

```

```
(Epoch 3 / 10) train acc: 0.450000; val_acc: 0.166000
(Iteration 7 / 20) loss: 1.809644
(Iteration 8 / 20) loss: 1.516281
(Epoch 4 / 10) train acc: 0.580000; val_acc: 0.189000
(Iteration 9 / 20) loss: 1.468933
(Iteration 10 / 20) loss: 1.123248
(Epoch 5 / 10) train acc: 0.590000; val_acc: 0.149000
(Iteration 11 / 20) loss: 1.259533
(Iteration 12 / 20) loss: 1.045123
(Epoch 6 / 10) train acc: 0.680000; val_acc: 0.172000
(Iteration 13 / 20) loss: 0.767402
(Iteration 14 / 20) loss: 0.886044
(Epoch 7 / 10) train acc: 0.790000; val_acc: 0.202000
(Iteration 15 / 20) loss: 0.858364
(Iteration 16 / 20) loss: 0.478033
(Epoch 8 / 10) train acc: 0.850000; val_acc: 0.200000
(Iteration 17 / 20) loss: 0.577705
(Iteration 18 / 20) loss: 0.403593
(Epoch 9 / 10) train acc: 0.870000; val_acc: 0.192000
(Iteration 19 / 20) loss: 0.667685
(Iteration 20 / 20) loss: 0.481448
(Epoch 10 / 10) train acc: 0.920000; val_acc: 0.188000
```

```
[5]: plt.subplot(2, 1, 1)
plt.plot(solver.loss_history, 'o')
plt.xlabel('iteration')
plt.ylabel('loss')

plt.subplot(2, 1, 2)
plt.plot(solver.train_acc_history, '-o')
plt.plot(solver.val_acc_history, '-o')
plt.legend(['train', 'val'], loc='upper left')
plt.xlabel('epoch')
plt.ylabel('accuracy')
plt.show()
```



## 1.2 Train the network

Now we train the 3 layer CNN on CIFAR-10 and assess its accuracy.

```
[6]: model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001)

solver = Solver(model, data,
                num_epochs=1, batch_size=50,
                update_rule='adam',
                optim_config={
                    'learning_rate': 1e-3,
                },
                verbose=True, print_every=20)
solver.train()
```

```
(Iteration 1 / 980) loss: 2.304400
(Epoch 0 / 1) train acc: 0.139000; val_acc: 0.137000
(Iteration 21 / 980) loss: 2.674096
(Iteration 41 / 980) loss: 1.943157
(Iteration 61 / 980) loss: 2.355880
```

(Iteration 81 / 980) loss: 1.891819  
(Iteration 101 / 980) loss: 2.012828  
(Iteration 121 / 980) loss: 1.873694  
(Iteration 141 / 980) loss: 1.755333  
(Iteration 161 / 980) loss: 1.682914  
(Iteration 181 / 980) loss: 1.815634  
(Iteration 201 / 980) loss: 1.682235  
(Iteration 221 / 980) loss: 2.011169  
(Iteration 241 / 980) loss: 1.638945  
(Iteration 261 / 980) loss: 1.787265  
(Iteration 281 / 980) loss: 1.933533  
(Iteration 301 / 980) loss: 1.681778  
(Iteration 321 / 980) loss: 1.808253  
(Iteration 341 / 980) loss: 1.513790  
(Iteration 361 / 980) loss: 1.446696  
(Iteration 381 / 980) loss: 1.577809  
(Iteration 401 / 980) loss: 1.748466  
(Iteration 421 / 980) loss: 1.770751  
(Iteration 441 / 980) loss: 1.728824  
(Iteration 461 / 980) loss: 1.434100  
(Iteration 481 / 980) loss: 1.824666  
(Iteration 501 / 980) loss: 1.649457  
(Iteration 521 / 980) loss: 1.504636  
(Iteration 541 / 980) loss: 1.610422  
(Iteration 561 / 980) loss: 1.462938  
(Iteration 581 / 980) loss: 1.568390  
(Iteration 601 / 980) loss: 1.557461  
(Iteration 621 / 980) loss: 2.082552  
(Iteration 641 / 980) loss: 1.487003  
(Iteration 661 / 980) loss: 1.418908  
(Iteration 681 / 980) loss: 1.800945  
(Iteration 701 / 980) loss: 1.792656  
(Iteration 721 / 980) loss: 1.611114  
(Iteration 741 / 980) loss: 1.502861  
(Iteration 761 / 980) loss: 1.696393  
(Iteration 781 / 980) loss: 1.589129  
(Iteration 801 / 980) loss: 1.463187  
(Iteration 821 / 980) loss: 1.648478  
(Iteration 841 / 980) loss: 1.932970  
(Iteration 861 / 980) loss: 1.610204  
(Iteration 881 / 980) loss: 1.471085  
(Iteration 901 / 980) loss: 1.505730  
(Iteration 921 / 980) loss: 1.184835  
(Iteration 941 / 980) loss: 1.605729  
(Iteration 961 / 980) loss: 1.466829  
(Epoch 1 / 1) train acc: 0.469000; val\_acc: 0.488000



## 2 Get > 65% validation accuracy on CIFAR-10.

In the last part of the assignment, we'll now ask you to train a CNN to get better than 65% validation accuracy on CIFAR-10.

### 2.0.1 Things you should try:

- Filter size: Above we used 7x7; but VGGNet and onwards showed stacks of 3x3 filters are good.
- Number of filters: Above we used 32 filters. Do more or fewer do better?
- Batch normalization: Try adding spatial batch normalization after convolution layers and vanilla batch normalization after affine layers. Do your networks train faster?
- Network architecture: Can a deeper CNN do better? Consider these architectures:
  - [conv-relu-pool]xN - conv - relu - [affine]xM - [softmax or SVM]
  - [conv-relu-pool]XN - [affine]XM - [softmax or SVM]
  - [conv-relu-conv-relu-pool]xN - [affine]xM - [softmax or SVM]

### 2.0.2 Tips for training

For each network architecture that you try, you should tune the learning rate and regularization strength. When doing this there are a couple important things to keep in mind:

- If the parameters are working well, you should see improvement within a few hundred iterations
- Remember the coarse-to-fine approach for hyperparameter tuning: start by testing a large range of hyperparameters for just a few training iterations to find the combinations of parameters that are working at all.
- Once you have found some sets of parameters that seem to work, search more finely around these parameters. You may need to train for more epochs.

```
[8]: # ===== #
# YOUR CODE HERE:
#   Implement a CNN to achieve greater than 65% validation accuracy
#   on CIFAR-10.
# ===== #

model = ThreeLayerConvNet( input_dim=(3, 32, 32), num_filters=32, filter_size=7,
                           hidden_dim=1000, num_classes=10, weight_scale=1e-3, reg=1e-3,
                           dtype=np.float32, use_batchnorm=False)
solver = Solver(model, data,
                num_epochs=10, batch_size=1000,
                update_rule='adam',
                lr_decay = 0.99,
                optim_config={
                    'learning_rate': 1e-3,
                },
                verbose=True, print_every=10)
solver.train()
```

```
# ===== #  
# END YOUR CODE HERE  
# ===== #
```

```
(Iteration 1 / 490) loss: 2.306692  
(Epoch 0 / 10) train acc: 0.119000; val_acc: 0.119000  
(Iteration 11 / 490) loss: 2.125276  
(Iteration 21 / 490) loss: 1.803713  
(Iteration 31 / 490) loss: 1.692730  
(Iteration 41 / 490) loss: 1.587318  
(Epoch 1 / 10) train acc: 0.477000; val_acc: 0.511000  
(Iteration 51 / 490) loss: 1.506039  
(Iteration 61 / 490) loss: 1.455957  
(Iteration 71 / 490) loss: 1.454727  
(Iteration 81 / 490) loss: 1.356845  
(Iteration 91 / 490) loss: 1.319754  
(Epoch 2 / 10) train acc: 0.559000; val_acc: 0.560000  
(Iteration 101 / 490) loss: 1.323770  
(Iteration 111 / 490) loss: 1.252332  
(Iteration 121 / 490) loss: 1.238275  
(Iteration 131 / 490) loss: 1.236292  
(Iteration 141 / 490) loss: 1.175816  
(Epoch 3 / 10) train acc: 0.601000; val_acc: 0.584000  
(Iteration 151 / 490) loss: 1.239669  
(Iteration 161 / 490) loss: 1.101122  
(Iteration 171 / 490) loss: 1.097917  
(Iteration 181 / 490) loss: 1.050101  
(Iteration 191 / 490) loss: 1.073609  
(Epoch 4 / 10) train acc: 0.653000; val_acc: 0.604000  
(Iteration 201 / 490) loss: 1.090483  
(Iteration 211 / 490) loss: 1.060792  
(Iteration 221 / 490) loss: 1.070918  
(Iteration 231 / 490) loss: 0.957499  
(Iteration 241 / 490) loss: 0.974188  
(Epoch 5 / 10) train acc: 0.694000; val_acc: 0.653000  
(Iteration 251 / 490) loss: 0.902507  
(Iteration 261 / 490) loss: 0.865422  
(Iteration 271 / 490) loss: 0.944982  
(Iteration 281 / 490) loss: 0.884322  
(Iteration 291 / 490) loss: 0.849965  
(Epoch 6 / 10) train acc: 0.723000; val_acc: 0.650000  
(Iteration 301 / 490) loss: 0.865890  
(Iteration 311 / 490) loss: 0.842320  
(Iteration 321 / 490) loss: 0.874008  
(Iteration 331 / 490) loss: 0.832901  
(Iteration 341 / 490) loss: 0.792152  
(Epoch 7 / 10) train acc: 0.761000; val_acc: 0.630000  
(Iteration 351 / 490) loss: 0.883844
```

```

(Iteration 361 / 490) loss: 0.781635
(Iteration 371 / 490) loss: 0.767589
(Iteration 381 / 490) loss: 0.802804
(Iteration 391 / 490) loss: 0.730505
(Epoch 8 / 10) train acc: 0.765000; val_acc: 0.653000
(Iteration 401 / 490) loss: 0.720166
(Iteration 411 / 490) loss: 0.714880
(Iteration 421 / 490) loss: 0.650941
(Iteration 431 / 490) loss: 0.713213
(Iteration 441 / 490) loss: 0.634855
(Epoch 9 / 10) train acc: 0.802000; val_acc: 0.686000
(Iteration 451 / 490) loss: 0.622079
(Iteration 461 / 490) loss: 0.646740
(Iteration 471 / 490) loss: 0.711556
(Iteration 481 / 490) loss: 0.690335
(Epoch 10 / 10) train acc: 0.823000; val_acc: 0.660000

```

```

[ ]: """
# I used here to see which parameters can be changed

model = ThreeLayerConvNet( input_dim=(3, 32, 32), num_filters=32, filter_size=7,
                           hidden_dim=100, num_classes=10, weight_scale=1e-3, reg=1e-3,
                           dtype=np.float32, use_batchnorm=True)

solver = Solver(model, data,
                num_epochs=10, batch_size=1000,
                update_rule='adam',
                lr_decay = 0.95,
                optim_config={
                    'learning_rate': 1e-3,
                },
                verbose=True, print_every=10)

solver.train()
"""

```

```

1  import numpy as np
2  from nndl.layers import *
3  import pdb
4
5
6  def conv_forward_naive(x, w, b, conv_param):
7      """
8      A naive implementation of the forward pass for a convolutional layer.
9
10     The input consists of N data points, each with C channels, height H and width
11     W. We convolve each input with F different filters, where each filter spans
12     all C channels and has height HH and width WW.
13
14     Input:
15     - x: Input data of shape (N, C, H, W)
16     - w: Filter weights of shape (F, C, HH, WW)
17     - b: Biases, of shape (F,)
18     - conv_param: A dictionary with the following keys:
19         - 'stride': The number of pixels between adjacent receptive fields in the
20           horizontal and vertical directions.
21         - 'pad': The number of pixels that will be used to zero-pad the input.
22
23     Returns a tuple of:
24     - out: Output data, of shape (N, F, H', W') where H' and W' are given by
25         H' = 1 + (H + 2 * pad - HH) / stride
26         W' = 1 + (W + 2 * pad - WW) / stride
27     - cache: (x, w, b, conv_param)
28     """
29     out = None
30     pad = conv_param['pad']
31     stride = conv_param['stride']
32
33     # ===== #
34     # YOUR CODE HERE:
35     # Implement the forward pass of a convolutional neural network.
36     # Store the output as 'out'.
37     # Hint: to pad the array, you can use the function np.pad.
38     # ===== #
39     N,C,H,W = x.shape
40     F,C,HH,WW = w.shape
41     H_out_shape = 1 + (H + 2 * pad - HH) // stride
42     W_out_shape = 1 + (W + 2 * pad - WW) // stride
43
44     out = np.zeros((N,F,H_out_shape,W_out_shape))
45
46     x = np.pad(x, pad_width = ((0,0),(0,0),(pad,pad),(pad,pad)), mode = 'constant')
47     for i in range(N):
48         for j in range(F):
49             for k in range(H_out_shape):
50                 for l in range(W_out_shape):
51                     x_selected = x[i,:,k * stride:(k*stride + HH), l * stride : (l * stride +
52                                     WW)]
53                     w_selected = w[j,:,:,:]
54                     out[i,j,k,l] = np.sum(x_selected * w_selected) + b[j]
55
56     # ===== #
57     # END YOUR CODE HERE
58     # ===== #
59
60     cache = (x, w, b, conv_param)
61     return out, cache
62
63 def conv_backward_naive(dout, cache):
64     """
65     A naive implementation of the backward pass for a convolutional layer.
66

```

```

67     Inputs:
68     - dout: Upstream derivatives.
69     - cache: A tuple of (x, w, b, conv_param) as in conv_forward_naive
70
71     Returns a tuple of:
72     - dx: Gradient with respect to x
73     - dw: Gradient with respect to w
74     - db: Gradient with respect to b
75     """
76     dx, dw, db = None, None, None
77
78     N, F, out_height, out_width = dout.shape
79     x, w, b, conv_param = cache
80
81     stride, pad = [conv_param['stride'], conv_param['pad']]
82     xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')
83     num_filts, _, f_height, f_width = w.shape
84
85     # ===== #
86     # YOUR CODE HERE:
87     # Implement the backward pass of a convolutional neural network.
88     # Calculate the gradients: dx, dw, and db.
89     # ===== #
90     N,F,H,W = x.shape
91
92     H_out_shape = 1 + (H - f_height) // stride
93     W_out_shape = 1 + (W - f_width) // stride
94
95     dx = np.zeros(x.shape)
96     dw = np.zeros(w.shape)
97     db = np.zeros(b.shape)
98
99     for i in range(N):
100         for j in range(num_filts):
101             if i == 0:
102                 db[j] += np.sum(dout[:,j,:,:])
103             for k in range(H_out_shape):
104                 for l in range(W_out_shape):
105                     k_tmp = k * stride
106                     l_tmp = l * stride
107                     dout_tmp = dout[i,j,k,l]
108                     dx[i,:, k_tmp:(k_tmp + f_height), l_tmp:(l_tmp + f_width)] += w[j,:,:, :]
109                     * dout_tmp
110                     dw[j,:,:, :] += x[i,:, k_tmp:(k_tmp + f_height), l_tmp:(l_tmp + f_width)]
111                     * dout_tmp
112                 dx = dx[:, :, pad:-pad, pad:-pad]
113             # ===== #
114             # END YOUR CODE HERE
115             # ===== #
116
117     return dx, dw, db
118
119 def max_pool_forward_naive(x, pool_param):
120     """
121     A naive implementation of the forward pass for a max pooling layer.
122
123     Inputs:
124     - x: Input data, of shape (N, C, H, W)
125     - pool_param: dictionary with the following keys:
126         - 'pool_height': The height of each pooling region
127         - 'pool_width': The width of each pooling region
128         - 'stride': The distance between adjacent pooling regions
129
130     Returns a tuple of:
131     - out: Output data
132     - cache: (x, pool_param)

```

```

132 """
133 out = None
134
135 # ===== #
136 # YOUR CODE HERE:
137 #   Implement the max pooling forward pass.
138 # ===== #
139
140 pool_height, pool_width, stride = pool_param['pool_height'], pool_param['pool_width'],
141 pool_param['stride']
142 N,C,H,W = x.shape
143
144 H_out_shape = 1 + (H - pool_height) // stride
145 W_out_shape = 1 + (W - pool_width) // stride
146 out = np.zeros((N,C, H_out_shape, W_out_shape))
147 for i in range(N):
148     for j in range(C):
149         for k in range(H_out_shape):
150             for l in range(W_out_shape):
151                 k_tmp = k * stride
152                 l_tmp = l * stride
153                 x_tmp = x[i,j,k_tmp:(k_tmp + pool_height),l_tmp:(l_tmp + pool_width)]
154                 out[i,j,k,l] = np.max(x_tmp)
155
156 # ===== #
157 # END YOUR CODE HERE
158 # ===== #
159 cache = (x, pool_param)
160 return out, cache
161
162 def max_pool_backward_naive(dout, cache):
163     """
164     A naive implementation of the backward pass for a max pooling layer.
165
166     Inputs:
167     - dout: Upstream derivatives
168     - cache: A tuple of (x, pool_param) as in the forward pass.
169
170     Returns:
171     - dx: Gradient with respect to x
172     """
173     dx = None
174     x, pool_param = cache
175     pool_height, pool_width, stride = pool_param['pool_height'], pool_param['pool_width'],
176     pool_param['stride']
177
178 # ===== #
179 # YOUR CODE HERE:
180 #   Implement the max pooling backward pass.
181 # ===== #
182
183 N,C,H,W = x.shape
184 H_out_shape = 1 + (H - pool_height) // stride
185 W_out_shape = 1 + (W - pool_width) // stride
186
187 dx = np.zeros((N,C,H,W))
188 for i in range(N):
189     for j in range(C):
190         for k in range(H_out_shape):
191             for l in range(W_out_shape):
192                 k_tmp = k * stride
193                 l_tmp = l * stride
194                 x_tmp = x[i,j,k_tmp:(k_tmp + pool_height),l_tmp:(l_tmp + pool_width)]
195                 dout_tmp = dout[i,j,k,l]
196                 din_mask = x_tmp == np.max(x_tmp)
197                 dx[i,j, k_tmp:(k_tmp + pool_height),l_tmp:(l_tmp + pool_width)] +=
198                 din_mask * dout_tmp

```

```

196 # ===== #
197 # END YOUR CODE HERE
198 # ===== #
199
200 return dx
201
202 def spatial_batchnorm_forward(x, gamma, beta, bn_param):
203     """
204     Computes the forward pass for spatial batch normalization.
205
206     Inputs:
207     - x: Input data of shape (N, C, H, W)
208     - gamma: Scale parameter, of shape (C,)
209     - beta: Shift parameter, of shape (C,)
210     - bn_param: Dictionary with the following keys:
211       - mode: 'train' or 'test'; required
212       - eps: Constant for numeric stability
213       - momentum: Constant for running mean / variance. momentum=0 means that
214         old information is discarded completely at every time step, while
215         momentum=1 means that new information is never incorporated. The
216         default of momentum=0.9 should work well in most situations.
217       - running_mean: Array of shape (D,) giving running mean of features
218       - running_var: Array of shape (D,) giving running variance of features
219
220     Returns a tuple of:
221     - out: Output data, of shape (N, C, H, W)
222     - cache: Values needed for the backward pass
223     """
224     out, cache = None, None
225
226     # ===== #
227     # YOUR CODE HERE:
228     #   Implement the spatial batchnorm forward pass.
229     #
230     #   You may find it useful to use the batchnorm forward pass you
231     #   implemented in HW #4.
232     # ===== #
233     N,C,H,W = x.shape
234     x_flattened = (x.reshape((N,H,W,C))).reshape((N*W*H,C))
235     out_bn, cache = batchnorm_forward(x_flattened, gamma,beta, bn_param = bn_param)
236     out = (out_bn.reshape((N,W,H,C))).swapaxes(1,3)
237     # ===== #
238     # END YOUR CODE HERE
239     # ===== #
240
241     return out, cache
242
243
244 def spatial_batchnorm_backward(dout, cache):
245     """
246     Computes the backward pass for spatial batch normalization.
247
248     Inputs:
249     - dout: Upstream derivatives, of shape (N, C, H, W)
250     - cache: Values from the forward pass
251
252     Returns a tuple of:
253     - dx: Gradient with respect to inputs, of shape (N, C, H, W)
254     - dgamma: Gradient with respect to scale parameter, of shape (C,)
255     - dbeta: Gradient with respect to shift parameter, of shape (C,)
256     """
257     dx, dgamma, dbeta = None, None, None
258
259     # ===== #
260     # YOUR CODE HERE:
261     #   Implement the spatial batchnorm backward pass.
262     #

```

```

263 # You may find it useful to use the batchnorm forward pass you
264 # implemented in HW #4.
265 # ===== #
266 N,C,H,W = dout.shape
267 dout_bn = dout.swapaxes(1,3).reshape((N*W*H,C))
268 dx_bn, dgamma_bn, dbeta_bn = batchnorm_backward(dout_bn,cache)
269 dx = dx_bn.reshape((N,C,H,W))
270 dgamma = dgamma_bn.reshape((C,))
271 dbeta = dbeta_bn.reshape((C,))
272 # ===== #
273 # END YOUR CODE HERE
274 # ===== #
275
276 return dx, dgamma, dbeta

```



```

1  import numpy as np
2
3  from nn1.layers import *
4  from nn1.conv_layers import *
5  from utils.fast_layers import *
6  from nn1.layer_utils import *
7  from nn1.conv_layer_utils import *
8
9  import pdb
10
11 class ThreeLayerConvNet(object):
12     """
13     A three-layer convolutional network with the following architecture:
14
15     conv - relu - 2x2 max pool - affine - relu - affine - softmax
16
17     The network operates on minibatches of data that have shape (N, C, H, W)
18     consisting of N images, each with height H and width W and with C input
19     channels.
20     """
21
22     def __init__(self, input_dim=(3, 32, 32), num_filters=32, filter_size=7,
23                 hidden_dim=100, num_classes=10, weight_scale=1e-3, reg=0.0,
24                 dtype=np.float32, use_batchnorm=False):
25         """
26         Initialize a new network.
27
28         Inputs:
29         - input_dim: Tuple (C, H, W) giving size of input data
30         - num_filters: Number of filters to use in the convolutional layer
31         - filter_size: Size of filters to use in the convolutional layer
32         - hidden_dim: Number of units to use in the fully-connected hidden layer
33         - num_classes: Number of scores to produce from the final affine layer.
34         - weight_scale: Scalar giving standard deviation for random initialization
35           of weights.
36         - reg: Scalar giving L2 regularization strength
37         - dtype: numpy datatype to use for computation.
38         """
39         self.use_batchnorm = use_batchnorm
40         self.params = {}
41         self.reg = reg
42         self.dtype = dtype
43
44
45         # ===== #
46         # YOUR CODE HERE:
47         #   Initialize the weights and biases of a three layer CNN. To initialize:
48         #   - the biases should be initialized to zeros.
49         #   - the weights should be initialized to a matrix with entries
50         #     drawn from a Gaussian distribution with zero mean and
51         #     standard deviation given by weight_scale.
52         # ===== #
53         C,H,W = input_dim
54         shapes = {}
55         shapes['W1'] = (num_filters, C, filter_size, filter_size)
56         shapes['W2'] = ((H//2) * (W//2) * num_filters, hidden_dim)
57         shapes['W3'] = (hidden_dim, num_classes)
58         shapes['b1'] = num_filters
59         shapes['b2'] = hidden_dim
60         shapes['b3'] = num_classes
61
62         for i in range(1,4):
63             str_W = 'W' + str(i)
64             str_b = 'b' + str(i)
65             self.params[str_W] = np.random.normal(loc = 0.0, scale = weight_scale, size =
66             shapes[str_W])
67             self.params[str_b] = np.zeros(shapes[str_b])

```

```

67
68 # ===== #
69 # END YOUR CODE HERE
70 # ===== #
71
72 for k, v in self.params.items():
73     self.params[k] = v.astype(dtype)
74
75
76 def loss(self, X, y=None):
77     """
78     Evaluate loss and gradient for the three-layer convolutional network.
79
80     Input / output: Same API as TwoLayerNet in fc_net.py.
81     """
82     W1, b1 = self.params['W1'], self.params['b1']
83     W2, b2 = self.params['W2'], self.params['b2']
84     W3, b3 = self.params['W3'], self.params['b3']
85
86     # pass conv_param to the forward pass for the convolutional layer
87     filter_size = W1.shape[2]
88     conv_param = {'stride': 1, 'pad': (filter_size - 1) / 2}
89
90     # pass pool_param to the forward pass for the max-pooling layer
91     pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}
92
93     scores = None
94
95     # ===== #
96     # YOUR CODE HERE:
97     # Implement the forward pass of the three layer CNN. Store the output
98     # scores as the variable "scores".
99     # ===== #
100
101     h1, cache1 = conv_relu_pool_forward(X, W1, b1, conv_param, pool_param)
102     h2, cache2 = affine_relu_forward(h1, W2, b2)
103     scores, cache3 = affine_forward(h2, W3, b3)
104
105     # ===== #
106     # END YOUR CODE HERE
107     # ===== #
108
109     if y is None:
110         return scores
111
112     loss, grads = 0, {}
113     # ===== #
114     # YOUR CODE HERE:
115     # Implement the backward pass of the three layer CNN. Store the grads
116     # in the grads dictionary, exactly as before (i.e., the gradient of
117     # self.params[k] will be grads[k]). Store the loss as "loss", and
118     # don't forget to add regularization on ALL weight matrices.
119     # ===== #
120
121     loss, dz = softmax_loss(scores, y)
122     loss += 0.5 * self.reg * (np.sum(W1**2) + np.sum(W2**2) + np.sum(W3**2))
123
124     dh2, dw3, grads['b3'] = affine_backward(dz, cache3)
125     dh1, dw2, grads['b2'] = affine_relu_backward(dh2, cache2)
126     _, dw1, grads['b1'] = conv_relu_pool_backward(dh1, cache1)
127
128     grads['W1'] = dw1 + self.reg * W1
129     grads['W2'] = dw2 + self.reg * W2
130     grads['W3'] = dw3 + self.reg * W3
131
132
133     # ===== #

```

```
134         # END YOUR CODE HERE
135         # ===== #
136
137         return loss, grads
138
139
140     pass
141
```

```

1  import numpy as np
2  import pdb
3
4
5  def affine_forward(x, w, b):
6      """
7      Computes the forward pass for an affine (fully-connected) layer.
8
9      The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
10     examples, where each example x[i] has shape (d_1, ..., d_k). We will
11     reshape each input into a vector of dimension D = d_1 * ... * d_k, and
12     then transform it to an output vector of dimension M.
13
14     Inputs:
15     - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
16     - w: A numpy array of weights, of shape (D, M)
17     - b: A numpy array of biases, of shape (M,)
18
19     Returns a tuple of:
20     - out: output, of shape (N, M)
21     - cache: (x, w, b)
22     """
23     out = None
24     # ===== #
25     # YOUR CODE HERE:
26     # Calculate the output of the forward pass. Notice the dimensions
27     # of w are D x M, which is the transpose of what we did in earlier
28     # assignments.
29     # ===== #
30     out = np.dot(x.reshape(x.shape[0], -1), w) + b
31
32
33     # ===== #
34     # END YOUR CODE HERE
35     # ===== #
36
37     cache = (x, w, b)
38     return out, cache
39
40
41  def affine_backward(dout, cache):
42      """
43      Computes the backward pass for an affine layer.
44
45      Inputs:
46      - dout: Upstream derivative, of shape (N, M)
47      - cache: Tuple of:
48        - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
49        - w: A numpy array of weights, of shape (D, M)
50        - b: A numpy array of biases, of shape (M,)
51
52      Returns a tuple of:
53      - dx: Gradient with respect to x, of shape (N, d_1, ..., d_k)
54      - dw: Gradient with respect to w, of shape (D, M)
55      - db: Gradient with respect to b, of shape (M,)
56      """
57      x, w, b = cache
58      dx, dw, db = None, None, None
59
60      # ===== #
61      # YOUR CODE HERE:
62      # Calculate the gradients for the backward pass.
63      # Notice:
64      #   dout is N x M
65      #   dx should be N x d_1 x ... x d_k; it relates to dout through multiplication with
66      #   w, which is D x M
67      #   dw should be D x M; it relates to dout through multiplication with x, which is N

```

```

67     x D after reshaping
68     # db should be M; it is just the sum over dout examples
69     # ===== #
69     flattened_x = x.reshape(x.shape[0],-1)
70     dx = np.dot(dout,w.T).reshape(x.shape)
71     dw = np.dot(flattened_x.T,dout)
72     db = np.sum(dout, axis = 0)
73
74     # ===== #
75     # END YOUR CODE HERE
76     # ===== #
77
78     return dx, dw, db
79
80 def relu_forward(x):
81     """
82     Computes the forward pass for a layer of rectified linear units (ReLU).
83
84     Input:
85     - x: Inputs, of any shape
86
87     Returns a tuple of:
88     - out: Output, of the same shape as x
89     - cache: x
90     """
91     # ===== #
92     # YOUR CODE HERE:
93     # Implement the ReLU forward pass.
94     # ===== #
95     relu = lambda x: x * (x > 0)
96     out = relu(x)
97     # ===== #
98     # END YOUR CODE HERE
99     # ===== #
100
101     cache = x
102     return out, cache
103
104
105 def relu_backward(dout, cache):
106     """
107     Computes the backward pass for a layer of rectified linear units (ReLU).
108
109     Input:
110     - dout: Upstream derivatives, of any shape
111     - cache: Input x, of same shape as dout
112
113     Returns:
114     - dx: Gradient with respect to x
115     """
116     x = cache
117
118     # ===== #
119     # YOUR CODE HERE:
120     # Implement the ReLU backward pass
121     # ===== #
122     dx = dout * (x > 0)
123
124     # ===== #
125     # END YOUR CODE HERE
126     # ===== #
127
128     return dx
129
130 def batchnorm_forward(x, gamma, beta, bn_param):
131     """
132     Forward pass for batch normalization.

```

During training the sample mean and (uncorrected) sample variance are computed from minibatch statistics and used to normalize the incoming data. During training we also keep an exponentially decaying running mean of the mean and variance of each feature, and these averages are used to normalize data at test-time.

At each timestep we update the running averages for mean and variance using an exponential decay based on the momentum parameter:

```
running_mean = momentum * running_mean + (1 - momentum) * sample_mean
running_var = momentum * running_var + (1 - momentum) * sample_var
```

Note that the batch normalization paper suggests a different test-time behavior: they compute sample mean and variance for each feature using a large number of training images rather than using a running average. For this implementation we have chosen to use running averages instead since they do not require an additional estimation step; the torch7 implementation of batch normalization also uses running averages.

Input:

- x: Data of shape (N, D)
- gamma: Scale parameter of shape (D,)
- beta: Shift parameter of shape (D,)
- bn\_param: Dictionary with the following keys:
  - mode: 'train' or 'test'; required
  - eps: Constant for numeric stability
  - momentum: Constant for running mean / variance.
  - running\_mean: Array of shape (D,) giving running mean of features
  - running\_var: Array of shape (D,) giving running variance of features

Returns a tuple of:

- out: of shape (N, D)
- cache: A tuple of values needed in the backward pass

```
mode = bn_param['mode']
eps = bn_param.get('eps', 1e-5)
momentum = bn_param.get('momentum', 0.9)
```

```
N, D = x.shape
running_mean = bn_param.get('running_mean', np.zeros(D, dtype=x.dtype))
running_var = bn_param.get('running_var', np.zeros(D, dtype=x.dtype))
```

```
out, cache = None, None
```

```
if mode == 'train':
```

```
    # ===== #
    # YOUR CODE HERE:
    #   A few steps here:
    #   (1) Calculate the running mean and variance of the minibatch.
    #   (2) Normalize the activations with the running mean and variance.
    #   (3) Scale and shift the normalized activations. Store this
    #       as the variable 'out'
    #   (4) Store any variables you may need for the backward pass in
    #       the 'cache' variable.
    # ===== #
    mean_x = np.mean(x, axis = 0)
    var_x = np.var(x, axis = 0)

    running_mean = momentum * running_mean + ( 1 - momentum ) * mean_x
    running_var = momentum * running_var + ( 1 - momentum ) * var_x

    standard_x = ( x - mean_x ) / ( np.sqrt(var_x + eps))

    out = gamma * standard_x + beta
    cache = (mean_x, var_x, standard_x, gamma, x, eps)
```

```

200         # ===== #
201         # END YOUR CODE HERE
202         # ===== #
203     elif mode == 'test':
204         # ===== #
205         # YOUR CODE HERE:
206         #     Calculate the testing time normalized activation. Normalize using
207         #     the running mean and variance, and then scale and shift appropriately.
208         #     Store the output as 'out'.
209         # ===== #
210
211         standard_x = ( x - running_mean) / (np.sqrt(running_var))
212         out = gamma * standard_x + beta
213         #cache = []
214         # ===== #
215         # END YOUR CODE HERE
216         # ===== #
217     else:
218         raise ValueError('Invalid forward batchnorm mode "%s"' % mode)
219
220     # Store the updated running means back into bn_param
221     bn_param['running_mean'] = running_mean
222     bn_param['running_var'] = running_var
223
224     return out, cache
225
226 def batchnorm_backward(dout, cache):
227     """
228     Backward pass for batch normalization.
229
230     For this implementation, you should write out a computation graph for
231     batch normalization on paper and propagate gradients backward through
232     intermediate nodes.
233
234     Inputs:
235     - dout: Upstream derivatives, of shape (N, D)
236     - cache: Variable of intermediates from batchnorm_forward.
237
238     Returns a tuple of:
239     - dx: Gradient with respect to inputs x, of shape (N, D)
240     - dgamma: Gradient with respect to scale parameter gamma, of shape (D,)
241     - dbeta: Gradient with respect to shift parameter beta, of shape (D,)
242     """
243     dx, dgamma, dbeta = None, None, None
244
245     # ===== #
246     # YOUR CODE HERE:
247     #     Implement the batchnorm backward pass, calculating dx, dgamma, and dbeta.
248     # ===== #
249     (mean_x, var_x, standard_x, gamma, x, eps) = cache
250     sample_size = x.shape[0]
251     sigma_x = np.sqrt(var_x + eps)
252
253     dgamma = np.sum(standard_x * dout,axis = 0)
254     dbeta = np.sum(dout, axis = 0)
255
256     dL_dx_st = dout * gamma
257
258     dx_st_da = 1 / sigma_x
259     dL_da = dx_st_da * dL_dx_st
260     da_dx = 1
261
262     dx_st_de = -0.5 * ( dx_st_da ** 3) * (x - mean_x)
263     dL_de = dx_st_de * dL_dx_st
264
265     dL_dvar = np.sum(dL_de, axis = 0)
266     dvar_dx = (2 * ( x - mean_x)) / sample_size

```

```

267
268 dL_dmean = np.sum(-dL_da, axis = 0)
269 dmean_dx = 1 / sample_size
270
271 dx = da_dx * dL_da + dvar_dx * dL_dvar + dmean_dx * dL_dmean
272 # ===== #
273 # END YOUR CODE HERE
274 # ===== #
275
276 return dx, dgamma, dbeta
277
278 def dropout_forward(x, dropout_param):
279     """
280     Performs the forward pass for (inverted) dropout.
281
282     Inputs:
283     - x: Input data, of any shape
284     - dropout_param: A dictionary with the following keys:
285       - p: Dropout parameter. We keep each neuron output with probability p.
286       - mode: 'test' or 'train'. If the mode is train, then perform dropout;
287         if the mode is test, then just return the input.
288       - seed: Seed for the random number generator. Passing seed makes this
289         function deterministic, which is needed for gradient checking but not in
290         real networks.
291
292     Outputs:
293     - out: Array of the same shape as x.
294     - cache: A tuple (dropout_param, mask). In training mode, mask is the dropout
295       mask that was used to multiply the input; in test mode, mask is None.
296     """
297     p, mode = dropout_param['p'], dropout_param['mode']
298     assert (0 < p <= 1), "Dropout probability is not in (0,1]"
299     if 'seed' in dropout_param:
300         np.random.seed(dropout_param['seed'])
301
302     mask = None
303     out = None
304
305     if mode == 'train':
306         # ===== #
307         # YOUR CODE HERE:
308         # Implement the inverted dropout forward pass during training time.
309         # Store the masked and scaled activations in out, and store the
310         # dropout mask as the variable mask.
311         # ===== #
312
313         mask = (np.random.rand(x.shape[0], x.shape[1]) < p) / p
314         out = x * mask
315         # ===== #
316         # END YOUR CODE HERE
317         # ===== #
318
319     elif mode == 'test':
320
321         # ===== #
322         # YOUR CODE HERE:
323         # Implement the inverted dropout forward pass during test time.
324         # ===== #
325
326         out = x
327
328         # ===== #
329         # END YOUR CODE HERE
330         # ===== #
331
332     cache = (dropout_param, mask)
333     out = out.astype(x.dtype, copy=False)

```



```

334
335     return out, cache
336
337 def dropout_backward(dout, cache):
338     """
339     Perform the backward pass for (inverted) dropout.
340
341     Inputs:
342     - dout: Upstream derivatives, of any shape
343     - cache: (dropout_param, mask) from dropout_forward.
344     """
345     dropout_param, mask = cache
346     mode = dropout_param['mode']
347
348     dx = None
349     if mode == 'train':
350         # ===== #
351         # YOUR CODE HERE:
352         # Implement the inverted dropout backward pass during training time.
353         # ===== #
354
355         dx = dout * mask
356
357         # ===== #
358         # END YOUR CODE HERE
359         # ===== #
360     elif mode == 'test':
361         # ===== #
362         # YOUR CODE HERE:
363         # Implement the inverted dropout backward pass during test time.
364         # ===== #
365
366         dx = dout
367
368         # ===== #
369         # END YOUR CODE HERE
370         # ===== #
371     return dx
372
373 def svm_loss(x, y):
374     """
375     Computes the loss and gradient using for multiclass SVM classification.
376
377     Inputs:
378     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
379         for the ith input.
380     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
381         0 <= y[i] < C
382
383     Returns a tuple of:
384     - loss: Scalar giving the loss
385     - dx: Gradient of the loss with respect to x
386     """
387     N = x.shape[0]
388     correct_class_scores = x[np.arange(N), y]
389     margins = np.maximum(0, x - correct_class_scores[:, np.newaxis] + 1.0)
390     margins[np.arange(N), y] = 0
391     loss = np.sum(margins) / N
392     num_pos = np.sum(margins > 0, axis=1)
393     dx = np.zeros_like(x)
394     dx[margins > 0] = 1
395     dx[np.arange(N), y] -= num_pos
396     dx /= N
397     return loss, dx
398
399
400 def softmax_loss(x, y):

```

```

401 """
402 Computes the loss and gradient for softmax classification.
403
404 Inputs:
405 - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
406     for the ith input.
407 - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
408     0 <= y[i] < C
409
410 Returns a tuple of:
411 - loss: Scalar giving the loss
412 - dx: Gradient of the loss with respect to x
413 """
414
415 probs = np.exp(x - np.max(x, axis=1, keepdims=True))
416 probs /= np.sum(probs, axis=1, keepdims=True)
417 N = x.shape[0]
418 loss = -np.sum(np.log(np.maximum(probs[np.arange(N), y], 1e-8))) / N
419 dx = probs.copy()
420 dx[np.arange(N), y] -= 1
421 dx /= N
422 return loss, dx
423

```

```

1  import numpy as np
2
3  """
4  This file implements various first-order update rules that are commonly used for
5  training neural networks. Each update rule accepts current weights and the
6  gradient of the loss with respect to those weights and produces the next set of
7  weights. Each update rule has the same interface:
8
9  def update(w, dw, config=None):
10
11  Inputs:
12      - w: A numpy array giving the current weights.
13      - dw: A numpy array of the same shape as w giving the gradient of the
14        loss with respect to w.
15      - config: A dictionary containing hyperparameter values such as learning rate,
16        momentum, etc. If the update rule requires caching values over many
17        iterations, then config will also hold these cached values.
18
19  Returns:
20      - next_w: The next point after the update.
21      - config: The config dictionary to be passed to the next iteration of the
22        update rule.
23
24  NOTE: For most update rules, the default learning rate will probably not perform
25  well; however the default values of the other hyperparameters should work well
26  for a variety of different problems.
27
28  For efficiency, update rules may perform in-place updates, mutating w and
29  setting next_w equal to w.
30  """
31
32
33  def sgd(w, dw, config=None):
34      """
35      Performs vanilla stochastic gradient descent.
36
37      config format:
38      - learning_rate: Scalar learning rate.
39      """
40      if config is None: config = {}
41      config.setdefault('learning_rate', 1e-2)
42
43      w -= config['learning_rate'] * dw
44      return w, config
45
46
47  def sgd_momentum(w, dw, config=None):
48      """
49      Performs stochastic gradient descent with momentum.
50
51      config format:
52      - learning_rate: Scalar learning rate.
53      - momentum: Scalar between 0 and 1 giving the momentum value.
54        Setting momentum = 0 reduces to sgd.
55      - velocity: A numpy array of the same shape as w and dw used to store a moving
56        average of the gradients.
57      """
58      if config is None: config = {}
59      config.setdefault('learning_rate', 1e-2)
60      config.setdefault('momentum', 0.9) # set momentum to 0.9 if it wasn't there
61      v = config.get('velocity', np.zeros_like(w)) # gets velocity, else sets it to zero.
62
63      # ===== #
64      # YOUR CODE HERE:
65      # Implement the momentum update formula. Return the updated weights
66      # as next_w, and the updated velocity as v.
67      # ===== #

```

```

68     v = config['momentum'] * v - config['learning_rate'] * dw
69     next_w = v + w
70     # ===== #
71     # END YOUR CODE HERE
72     # ===== #
73
74     config['velocity'] = v
75
76     return next_w, config
77
78 def sgd_nesterov_momentum(w, dw, config=None):
79     """
80     Performs stochastic gradient descent with Nesterov momentum.
81
82     config format:
83     - learning_rate: Scalar learning rate.
84     - momentum: Scalar between 0 and 1 giving the momentum value.
85       Setting momentum = 0 reduces to sgd.
86     - velocity: A numpy array of the same shape as w and dw used to store a moving
87       average of the gradients.
88     """
89     if config is None: config = {}
90     config.setdefault('learning_rate', 1e-2)
91     config.setdefault('momentum', 0.9) # set momentum to 0.9 if it wasn't there
92     v = config.get('velocity', np.zeros_like(w)) # gets velocity, else sets it to zero.
93
94     # ===== #
95     # YOUR CODE HERE:
96     #   Implement the momentum update formula. Return the updated weights
97     #   as next_w, and the updated velocity as v.
98     # ===== #
99     v_prev = v
100    v = config['momentum'] * v_prev - config['learning_rate'] * dw
101    next_w = v + w + config['momentum'] * (v - v_prev)
102    # ===== #
103    # END YOUR CODE HERE
104    # ===== #
105
106    config['velocity'] = v
107
108    return next_w, config
109
110 def rmsprop(w, dw, config=None):
111     """
112     Uses the RMSProp update rule, which uses a moving average of squared gradient
113     values to set adaptive per-parameter learning rates.
114
115     config format:
116     - learning_rate: Scalar learning rate.
117     - decay_rate: Scalar between 0 and 1 giving the decay rate for the squared
118       gradient cache.
119     - epsilon: Small scalar used for smoothing to avoid dividing by zero.
120     - beta: Moving average of second moments of gradients.
121     """
122     if config is None: config = {}
123     config.setdefault('learning_rate', 1e-2)
124     config.setdefault('decay_rate', 0.99)
125     config.setdefault('epsilon', 1e-8)
126     config.setdefault('a', np.zeros_like(w))
127
128     next_w = None
129
130     # ===== #
131     # YOUR CODE HERE:
132     #   Implement RMSProp. Store the next value of w as next_w. You need
133     #   to also store in config['a'] the moving average of the second
134     #   moment gradients, so they can be used for future gradients. Concretely,

```

```

135 # config['a'] corresponds to "a" in the lecture notes.
136 # ===== #
137 config['a'] = config['decay_rate'] * config['a'] + (1 - config['decay_rate']) * dw *
dw
138 c = 1 / (np.sqrt(config['a']) + config['epsilon'])
139 next_w = w - config['learning_rate'] * c * dw
140 # ===== #
141 # END YOUR CODE HERE
142 # ===== #
143
144 return next_w, config
145
146
147 def adam(w, dw, config=None):
148     """
149     Uses the Adam update rule, which incorporates moving averages of both the
150     gradient and its square and a bias correction term.
151
152     config format:
153     - learning_rate: Scalar learning rate.
154     - beta1: Decay rate for moving average of first moment of gradient.
155     - beta2: Decay rate for moving average of second moment of gradient.
156     - epsilon: Small scalar used for smoothing to avoid dividing by zero.
157     - m: Moving average of gradient.
158     - v: Moving average of squared gradient.
159     - t: Iteration number.
160     """
161     if config is None: config = {}
162     config.setdefault('learning_rate', 1e-3)
163     config.setdefault('beta1', 0.9)
164     config.setdefault('beta2', 0.999)
165     config.setdefault('epsilon', 1e-8)
166     config.setdefault('v', np.zeros_like(w))
167     config.setdefault('a', np.zeros_like(w))
168     config.setdefault('t', 0)
169
170     next_w = None
171
172     # ===== #
173     # YOUR CODE HERE:
174     # Implement Adam. Store the next value of w as next_w. You need
175     # to also store in config['a'] the moving average of the second
176     # moment gradients, and in config['v'] the moving average of the
177     # first moments. Finally, store in config['t'] the increasing time.
178     # ===== #
179
180     config['t'] += 1
181     config['v'] = config['beta1'] * config['v'] + (1 - config['beta1']) * dw
182     config['a'] = config['beta2'] * config['a'] + (1 - config['beta2']) * dw * dw
183
184     v_tld = config['v'] / (1 - config['beta1'] ** config['t'])
185     a_tld = config['a'] / (1 - config['beta2'] ** config['t'])
186
187     c = 1 / (np.sqrt(a_tld) + config['epsilon'])
188     next_w = w - config['learning_rate'] * v_tld * c
189
190     # ===== #
191     # END YOUR CODE HERE
192     # ===== #
193
194     return next_w, config
195

```